

Report for Project 3: Collaboration and Competition

Udacity Deep Reinforcement Learning Nanodegree

April 27, 2020

1 The Implementation

The main part of the code is written in the Jupyter Notebook `Tennis.ipynb`, which guides the user when executing the code. The intelligent agent has been implemented in the `agent.py` file. This Multi-Agent DDPG (MADDPG) agent consists of two DDPG agents defined in `ddpg.py`, each of which uses two neural networks, for its actor and its critic, as defined in `model.py`. The prioritized experience replay buffer is defined in `buffer.py`, while the Ornstein-Uhlenbeck noise process is defined in `OUnoise.py`.

1.1 The Neural Networks

The networks as defined in `model.py` by default have two hidden fully-connected layers with 256 and 128 neurons, respectively. The goal of the actor is to provide the best action given a state s : $a = \mu(s; \theta_\mu)$. The critic's goal is to estimate Q for the best-believed action, i.e. $Q(s, \mu(s; \theta_\mu); \theta_Q)$. For the critic, the actions are fed into the network as inputs to the second hidden layer, while states are directly input to the first layer. Here, always the whole state space for both agents is used. For the critic, also two action sets are used, namely

- own actions and own next actions for the critic update and experience weights for the Replay Buffer
- own actions and expected actions for the actor update.

Important to notice here is the fact that the order of these two sets had to be swapped for the two agents, since they are playing against each other (see the `if self.index:` decision in `ddpg.learn`).

Replay Buffer	Value	Agent Hypers	Value
Buffer Size	100k	γ	0.99
Batch Size	128	τ	0.06
α	0	$\text{LRate}_{\text{actor}}$	0.001
β	$0 \rightarrow 1$	$\text{LRate}_{\text{critic}}$	0.001

Table 1: The hyperparameters used for the learning algorithm.

1.2 The DDPG Class

The agent class is capable of learning according to the Deep Deterministic Policy Gradient (DDPG) algorithm as discussed in class and in this paper.

Within its `__init__` method, several internal variables are defined: state size, action size, and the `PrioritizedReplayBuffer` for experience replay. Besides, the network weights (local w and target w^-) and optimizer to be used (Adam) are defined for the agent’s actor and critic network. As action noise, an Ornstein-Uhlenbeck process is defined.

The `step` method adds the current experience to the replay buffer and triggers a learning step. To do this, it triggers the `learn` method, which backpropagates the errors made from sampled experiences through the network to update weights, once for the critic, and once for the actor. The actual (soft) update is done in `soft_update`. Since we are using experience replay, when adding samples in the `step` function, we have to calculate the importance weights (priorities) for the samples. For us, this is the quadratic error between expected Q (based on local critic) and target Q (based on target critic). Similarly, when learning from samples, we need to update the weights (priorities) for the sampled tuples, by checking their current expected Q value against the current target Q value (again, quadratic error).

Hyperparameters used are mentioned in Table 1.

The `act` method is used to return the action suggested by the actor for the current state.

The `save` and `load_weights` functions are used to save trained agents for further use and to load previously-trained agents, respectively. They are quite straightforward, only some checks for stringent dimensions of the networks are done in the `load` method.

The `PrioritizedReplayBuffer` class was based on this paper and has been inspired by PHRABAL’s implementation. It comprises an initialization method as well as methods for adding experiences and corresponding weights, sampling from the stored experiences, updating priorities, and a function for returning the current length of the buffer. The value β , which controls the importance of sampling weights, is increased linearly in the training loop from 0 to 1.




Figure 1: The scores reached by the intelligent agent.

2 Results

As Figure 1 shows, the agent is capable of learning, the environment has been solved within **711 episodes**.

3 Experiences and Insights

This project was quite challenging. Most notably, it the learning is facing quite big instabilities in this environment. Due to the very time-consuming runs (even in the GPU-enabled Udacity work space), it is hard to find suitable hyperparameters. After trying different model sizes, buffer sizes, with and without experience replay, and with several configurations of the models (completely independent regarding states or not, with or without normalization, with our without dropouts, etc.), I finally got a hold on the environment. Both DDPG agents share now state spaces for the actor (input size = $2 \times \text{state size}$), while they use own and next actions for the critic.

4 Ideas for Future Improvements

While the setup currently works sufficiently well, it may be interesting to re-implement some of the features tried on the way which have not found their way into the final solving version. This could be

- Adding dropouts to the critic
- Adjust hyperparameters, especially the batch size and prioritization parameters.

Additionally, a mentor suggested to use parameter noise instead of action space noise as described here (video here).

It would be interesting to find a better configuration which potentially could solve the environment even faster.